

Kernel Pwn Cheat Sheet

- [Kernel version](#)
- [Kernel config](#)
- [Process management](#)
 - [task_struct](#)
 - [current](#)
- [Syscall](#)
- [Memory allocator](#)
 - [kmem_cache](#)
 - [kmalloc](#)
 - [kfree](#)
- [Physmem](#)
- [Snippet](#)
- [Structures](#)
 - [ldt_struct](#)
 - [shm_file_data](#)
 - [seq_operations](#)
 - [msg_msg, msg_msgseg](#)
 - [subprocess_info](#)
 - [timerfd_ctx](#)
 - [pipe_buffer](#)
 - [tty_struct](#)
 - [setxattr](#)
 - [sk_buff](#)
- [Variables](#)
 - [modprobe_path](#)
 - [core_pattern](#)
 - [n_tty_ops](#)

Kernel version

```
commit 09688c0166e76ce2fb85e86b9d99be8b0084cdf9 (HEAD -> master, t
Author: Linus Torvalds <torvalds@linux-foundation.org>
Date:   Sun Mar 13 13:23:37 2022 -0700
```

Linux 5.17-rc8

Kernel config

config	memo
CONFIG_KALLSYMS	/proc/sys/kernel/kptr_restrict
CONFIG_USERFAULTFD	/proc/sys/vm/unprivileged_use
CONFIG_STATIC_USERMODEHELPER	
CONFIG_SLUB	default allocator
CONFIG_SLAB	
CONFIG_SLAB_FREELIST_RANDOM	
CONFIG_SLAB_FREELIST_HARDENED	
CONFIG_FG_KASLR	
CONFIG_BPF	/proc/sys/kernel/unprivileged_l
CONFIG_SMP	multi-processor

Process management

task_struct

- [task_struct](#)
 - [thread_info](#)
 - `syscall_work`
 - [cred](#)
 - `tasks`

- `init_task`
 - `init_cred`
- `comm`
 - `prctl(PR_SET_NAME, name);`

current

- `current`
 - `get_current`
 - `current_task`
 - `DECLARE_PER_CPU`
 - `DECLARE_PER_CPU_SECTION`
 - `__PCPU_ATTRS`
 - *case CONFIG_SMP*
 - `PER_CPU_BASE_SECTION`
 - `this_cpu_read_stable`
 - `__pcpu_size_call_return`
 - `this_cpu_read_stable_8`
 - `percpu_stable_op`
 - *case CONFIG_SMP*
 - `movq %%gs:%P[var], %[val]` where
`var = ¤t_task`
- `start_kernel`
 - `setup_per_cpu_areas`
 - *case CONFIG_SMP*
 - `per_cpu_offset`
 - `__per_cpu_offset[cpu] = pcpu_base_addr -
__per_cpu_start + pcpu_unit_offsets[cpu]`
 - `switch_to_new_gdt`
 - `load_percpu_segment`
 - `cpu_kernelmode_gs_base`
 - `fixed_percpu_data`
 - `DECLARE_PER_CPU_FIRST`
 - `fixed_percpu_data`
 - `per_cpu`
 - *case CONFIG_SMP*

- `per_cpu_ptr`
 - `SHIFT_PERCPU_PTR`
 - `RELOC_HIDE`
- `case CONFIG_SMP`
 - `gs = &fixed_percpu_data.gs_base +`
`__per_cpu_offset[cpu]`

Syscall

- `entry_SYSCALL_64`
 - `pt_regs`
 - `pt_regs` may be use for stack pivoting
 - `do_syscall_64`
 - `add_random_kstack_offset()`;
 - `syscall_enter_from_user_mode`
 - `__syscall_enter_from_user_work`
 - `syscall_trace_enter`
 - `SYSCALL_WORK_SECCOMP`
 - `do_syscall_x64`
 - `swapgs_restore_regs_and_return_to_usermode`

Memory allocator

kmem_cache

- `case CONFIG_SLUB`
 - `kmem_cache`
 - `kmem_cache_cpu`
 - `freelist`
 - `slab`
 - `slab_cache`
 - `freelist`
 - `offset`
 - `random`

- [kmem_cache_node](#)
- *case CONFIG_SLAB*
 - [kmem_cache](#)
 - [array_cache](#)
 - `entry`
 - [kmem_cache_node](#)
 - `shared`

kmalloc

- [kmalloc](#)
 - [kmalloc_index](#)
 - [__kmalloc_index](#)
 - *case CONFIG_SLUB*
 - `#define KMALLOCC_MIN_SIZE 8`
 - *case CONFIG_SLAB*
 - `#define KMALLOCC_MIN_SIZE 32`
 - [kmalloc_caches](#)
 - [kmalloc_type](#)
 - `#define GFP_KERNEL_ACCOUNT (GFP_KERNEL | __GFP_ACCOUNT)`
 - `GFP_KERNEL → KMALLOCC_NORMAL`
 - `GFP_KERNEL_ACCOUNT → KMALLOCC_CGROUP`
 - *case CONFIG_SLUB*
 - [kmem_cache_alloc_trace](#)
 - [slab_alloc](#)
 - [slab_alloc_node](#)
 - [__slab_alloc](#)
 - [__slab_alloc](#)
 - [new_slab](#)
 - [allocate_slab](#)
 - [shuffle_freelist](#)
 - [get_freepointer_safe](#)
 - [freelist_ptr](#)

- `get_freepointer_safe(cache, object) = (object + cache->offset) ^ *(object + cache->offset) ^ cache->random`
- *case CONFIG_SLAB*
 - `kmem_cache_alloc_trace`
 - `slab_alloc`
 - `__do_cache_alloc`
 - `____cache_alloc`
 - `cache_alloc_refill`
 - `____cache_alloc_node`
 - `cache_grow_begin`
 - `cache_init_objs`
 - `shuffle_freelist`

kfree

- *case CONFIG_SLUB*
 - `kfree`
 - `slab_free`
 - `do_slab_free`
 - `likely(slab == c->slab) → likely(slab == slab->slab_cache->cpu_slab->slab)`
 - `__slab_free`
 - `set_freepointer`
 - `BUG_ON(object == fp);`
- *case CONFIG_SLAB*
 - `kfree`
 - `____cache_free`
 - `cache_flusharray`
 - `__free_one`
 - `WARN_ON_ONCE(ac->avail > 0 && ac->entry[ac->avail - 1] == objp)`

Physem

- `page tables`

- `page_offset_base`
 - heap base address (by `kmalloc`) and it is mapped to `/dev/mem`
 - `secondary_startup_64` can be found at `page_offset_base + offset`
- `vmalloc_base`
- `vmemmap_base`
- [page](#)
- [vmalloc_to_page](#)
- [page_to_virt](#)
 - `page_to_virt(page) = page_offset_base + (((page - vmemmap_base) / 64) << 12)`
 - [__va](#)
 - [PAGE_OFFSET](#)
 - [__PAGE_OFFSET](#)
 - `page_offset_base`
 - [PFN_PHYS](#)
 - [PAGE_SHIFT](#)
 - [page_to_pfn](#)
 - [__page_to_pfn](#)
 - [vmemmap](#)
 - [VMEMMAP_START](#)
 - `vmemmap_base`

Snippet

- gain root privileges
 - (kernel) `commit_creds(prepare_kernel_cred(NULL));`
- break out of namespaces
 - (kernel) `switch_task_namespaces(find_task_by_vpid(1), init_nsproxy);`
 - (user) `setns(open("/proc/1/ns/mnt", O_RDONLY), 0);`
 - (user) `setns(open("/proc/1/ns/pid", O_RDONLY), 0);`
 - (user) `setns(open("/proc/1/ns/net", O_RDONLY), 0);`

Structures

| structure | size | flag (v5.14+) | memo |
|-----------------|---------------------|--------------------|-------------------------------|
| ldt_struct | 16 | GFP_KERNEL_ACCOUNT | |
| shm_file_data | 32 | GFP_KERNEL | |
| seq_operations | 32 | GFP_KERNEL_ACCOUNT | /proc/self/stat |
| msg_msg | 48 ~
4096 | GFP_KERNEL_ACCOUNT | |
| msg_msgseg | 8 ~
4096 | GFP_KERNEL_ACCOUNT | |
| subprocess_info | 96 | GFP_KERNEL | socket(22,
AF_INET,
0); |
| timerfd_ctx | 216 | GFP_KERNEL | |
| pipe_buffer | 640
= 40
x 16 | GFP_KERNEL_ACCOUNT | |
| tty_struct | 696 | GFP_KERNEL | /dev/ptmx |
| setxattr | 0 ~ | GFP_KERNEL | |
| sk_buff | 320
~ | GFP_KERNEL_ACCOUNT | |

ldt_struct

- `modify_ldt`
 - `write_ldt`
 - `alloc_ldt_struct`
 - `read_ldt`
 - `desc_struct`
 - `copy_to_user`

- `copy_to_user` won't panic the kernel when accessing wrong address

shm_file_data

- `shmat`
 - `do_shmat`

seq_operations

- `proc_stat_init`
 - `stat_proc_ops`
- `stat_open`
 - `single_open_size`
 - `single_open`
- `seq_read_iter`
 - `m->op->start`

msg_msg, msg_msgseg

- `msgsnd`
 - `ksys_msgsnd`
 - `do_msgsnd`
 - `load_msg`
 - `alloc_msg`
- `msgrcv`
 - `ksys_msgrcv`
 - `do_msgrcv`
 - `#define MSG_COPY 040000`

subprocess_info

- `socket`
 - `__sys_socket`
 - `sock_create`
 - `__sock_create`
 - `__request_module`

- `call_modprobe`
 - `call_usermodehelper_setup`

timerfd_ctx

- `timerfd_create`
- `timerfd_release`
 - `kfree_rcu`

pipe_buffer

- `pipe`, `pipe2`
 - `do_pipe2`
 - `do_pipe_flags`
 - `create_pipe_files`
 - `get_pipe_inode`
 - `alloc_pipe_info`
 - `#define PIPE_DEF_BUFFERS 16`
 - `pipefifo_fops`
- `pipe_write`
 - `buf->ops = &anon_pipe_buf_ops;`
- `pipe_release`
 - `put_pipe_info`
 - `free_pipe_info`
 - `pipe_buf_release`
 - `ops->release`

tty_struct

- `unix98_pty_init`
 - `tty_default_fops`
 - `tty_fops`
- `ptmx_open`
 - `tty_init_dev`
 - `alloc_tty_struct`
- `tty_ioctl`
 - `tty_paranoia_check`

- `#define TTY_MAGIC 0x5401`
- `tty_pair_get_tty`
- `tty->ops->iocctl`

setxattr

- `setxattr`
 - `path_setxattr`
 - `setxattr`
 - `vfs_setxattr` may fail, but `kvmalloc` and `kvmfree` complete successfully

sk_buff

- `socketpair`
 - `__sys_socketpair`
 - `sock_create`
 - `__sock_create`
 - `case PF_UNIX`
 - `unix_family_ops`
 - `unix_create`
 - `case SOCK_DGRAM`
 - `unix_dgram_ops`
 - `unix_create1`
 - `sk->sk_allocation = GFP_KERNEL_ACCOUNT;`
 - `unix_dgram_sendmsg`
 - `sock_alloc_send_skb`
 - `alloc_skb_with_frags`
 - `alloc_skb`
 - `__alloc_skb`
 - `struct skb_shared_info` is at the end of `data`

Variables

| variable | memo |
|----------|------|
| | |

| | |
|---------------|--|
| modprobe_path | /proc/sys/kernel/modprobe |
| core_pattern | /proc/sys/kernel/core_pattern |
| n_tty_ops | (read) <code>scanf</code> , (ioctl) <code>fgets</code> |

modprobe_path

- `execve`
 - `do_execve`
 - `do_execveat_common`
 - `bprm_execve`
 - `exec_binprm`
 - `search_binary_handler`
 - `__request_module`
 - `call_modprobe`
 - `call_usermodehelper_setup`
 - `call_usermodehelper_exec`

core_pattern

- `do_coredump`
 - `format_corename`
 - `call_usermodehelper_setup`
 - `call_usermodehelper_exec`

n_tty_ops

- `tty_struct`
 - `tty_ldisc`
- `n_tty_init`
 - `tty_register_ldisc`